

Michael J. Paul

michael.j.paul@colorado.edu

<http://cs.jhu.edu/~mpaul/>

Last Updated: September 23, 2015

Research interests: natural language processing (especially for the web and social media), machine learning (especially unsupervised and semi-supervised), (probabilistic) topic modeling, health and medical informatics, computational epidemiology, computational social science

EMPLOYMENT

2015 – Present	University of Colorado Boulder <i>Assistant Professor, Information Science</i> <i>Assistant Professor by Courtesy, Computer Science</i>
----------------	--

EDUCATION

2010 – 2015	Johns Hopkins University <i>Ph.D.</i> , Computer Science, August 2015 <i>M.S.E.</i> , Computer Science, May 2012 <ul style="list-style-type: none">Advisors: Mark Dredze (primary) and Jason EisnerCommittee: Mark Dredze, Jason Eisner, Eric Horvitz, Hanna Wallach
2005 – 2009	University of Illinois at Urbana-Champaign <i>B.S. with Honors</i> , Computer Science, May 2009
Spring 2008	Lund University Study abroad: Lund, Sweden

AWARDS AND HONORS

2013 – 2015	Microsoft Research Fellowship	Microsoft Corporation
	<ul style="list-style-type: none">One of 12 recipients across North America, and the only recipient in natural language processing	
2010 – 2012	Graduate Research Fellowship	National Science Foundation
2010 – 2015	Dean's Fellowship	Johns Hopkins University
2005 – 2008	Edmund J. James Scholar	University of Illinois
	<ul style="list-style-type: none">Undergraduate honors program	

TEACHING

12/04/2013	Guest Lecture	Johns Hopkins University
	<ul style="list-style-type: none">CS465: Natural Language ProcessingInstructor: Dr. Jason EisnerTitle: An introduction to topic models	
10/30/2013	Guest Lecture	University of Pennsylvania
	<ul style="list-style-type: none">NETS 213: Crowdsourcing & Human ComputationInstructor: Dr. Chris Callison-BurchTitle: Mining health information from TwitterMy slides were reused for Fall 2014	

Spring 2013	Teaching Assistant <ul style="list-style-type: none"> • CS476/676: Machine Learning in Complex Domains • Instructor: Dr. Suchi Saria • First time this course was offered 	Johns Hopkins University
Fall 2012	Teaching Assistant <ul style="list-style-type: none"> • CS475: Machine Learning • Instructor: Dr. Mark Dredze • Lectured on: 9/26/2012 (Generative vs discriminative learning), 10/8/2012 (Kernel methods) • Substitute lecturer: 2/10/2014 (Regression) 	Johns Hopkins University

JOURNAL PUBLICATIONS

2015	<p>David A. Broniatowski, Mark Dredze, Michael J. Paul, Andrea Dugas. Using social media to perform local influenza surveillance in an inner-city hospital: a retrospective observational study. <i>JMIR Public Health and Surveillance</i> 1(1): e5.</p> <p>Michael J. Paul and Mark Dredze. SPRITE: Generalizing topic models with structured priors. <i>Transactions of the Association for Computational Linguistics (TACL)</i> 3: 43-57.</p> <p>Shiliang Wang, Michael J. Paul, Mark Dredze. Social media as a sensor of air quality and public response in China. <i>Journal of Medical Internet Research (JMIR)</i> 17(3): e22.</p>	
2014	<p>David A. Broniatowski, Michael J. Paul, Mark Dredze. Twitter: Big data opportunities (letter). <i>Science</i> 345(6193): 148.</p> <p>Michael J. Paul, Mark Dredze, David Broniatowski. Twitter improves influenza forecasting. <i>PLOS Currents Outbreaks</i>. October 2014. doi: 10.1371/currents.outbreaks.90b9ed0f59bae4ccaa683a39865d9117.</p> <p>Michael J. Paul and Mark Dredze. Discovering health topics in social media using topic models. <i>PLOS ONE</i> 9(8): e103408.</p> <p>Ahmed Abbasi, Donald Adjeroh, Mark Dredze, Michael J. Paul, Fatemeh Mariam Zahedi, Huimin Zhao, Nitin Walia, Hemant Jain, Patrick Sanvanson, Reza Shaker, Marco D. Huesch, Richard Beal, Wanhong Zheng, Marie Abate, Arun Ross. Social media analytics for smart health. <i>IEEE Intelligent Systems</i> 29(2):60-80.</p> <p>Byron C. Wallace, Michael J. Paul, Urmimala Sarkar, Thomas A. Trikalinos, Mark Dredze. A large-scale quantitative analysis of latent factors and sentiment in online doctor reviews. <i>Journal of the American Medical Informatics Association (JAMIA)</i> 21(6), 1098-1103.</p>	
2013	<p>David A. Broniatowski, Michael J. Paul, Mark Dredze. National and local influenza surveillance through Twitter: an analysis of the 2012-2013 influenza epidemic. <i>PLOS ONE</i> 8(12): e83672.</p>	
2011	<p>Roxana Girju and Michael J. Paul. Modeling reciprocity in social interactions with probabilistic latent space models. <i>Natural Language Engineering</i> 17(1). Cambridge University Press 2011.</p>	

PEER-REVIEWED CONFERENCE PUBLICATIONS

- 2015 Michael J. Paul, Ryen W. White, Eric Horvitz. **Diagnoses, decisions, and outcomes: Web search as medical decision support for cancer.** *24th International World Wide Web Conference (WWW 2015)*, Florence, Italy. May 2015.
- 2013 Alex Lamb, Michael J. Paul, Mark Dredze. **Separating fact from fear: Tracking flu infections on Twitter.** *2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT 2013)*, Atlanta. June 2013.
- Michael J. Paul and Mark Dredze. **Drug extraction from the Web: Summarizing drug experiences with multi-dimensional topic models.** *2013 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT 2013)*, Atlanta. June 2013.
- 2012 Michael J. Paul and Mark Dredze. **Factorial LDA: Sparse multi-dimensional models of text.** In *Advances in Neural Information Processing Systems (NIPS)*. Lake Tahoe, Nevada. December 2012.
- Michael J. Paul. **Mixed membership Markov models for unsupervised conversation modeling.** In the *2012 Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL 2012)*, Jeju, Korea. July 2012.
- Michael J. Paul and Jason Eisner. **Implicitly intersecting weighted automata using dual decomposition.** In the *2012 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT 2012)*, Montreal, Canada. June 2012.
- 2011 Michael J. Paul and Mark Dredze. **You are what you Tweet: Analyzing Twitter for public health.** In the *5th International AAAI Conference on Weblogs and Social Media (ICWSM 2011)*. July 2011.
- Delip Rao, Michael Paul, Clayton Fink, David Yarowsky, Timothy Oates, Glen Coppersmith. **Hierarchical Bayesian models for latent attribute detection in social networks.** In the *5th International AAAI Conference on Weblogs and Social Media (ICWSM 2011)*. July 2011.
- 2010 Michael J. Paul, ChengXiang Zhai and Roxana Girju. **Summarizing contrastive viewpoints in opinionated text.** In the proceedings of the *2010 Conference on Empirical Methods in Natural Language Processing (EMNLP 2010)*, Cambridge, Massachusetts. October 2010.
- Michael Paul and Roxana Girju. **Comparative scientific research analysis with a language-independent cross-collection model.** In the proceedings of *XXV Congreso de la Sociedad Española para el Procesamiento del Lenguaje Natural (SEPLN-2010)*, Valencia, Spain, September 2010.
- Michael Paul and Roxana Girju. **A two-dimensional topic-aspect model for discovering multi-faceted topics.** In the proceedings of the *Twenty-Fourth AAAI Conference on Artificial Intelligence (AAAI 2010)*, Atlanta, Georgia, July 2010.
- 2009 Michael Paul and Roxana Girju. **Topic modeling of research fields: an interdisciplinary perspective.** In the proceedings of *Recent Advances in Natural Language Processing (RANLP 2009)*. September 2009.

Michael Paul and Roxana Girju. **Cross-cultural analysis of blogs and forums with mixed-collection topic models.** *Empirical Methods in Natural Language Processing (EMNLP 2009)*, Singapore. August 2009.

Michael Paul, Roxana Girju, Chen Li. **Mining the Web for reciprocal relationships.** In the proceedings of the *13th Conference on Computational Natural Language Learning (CoNLL 2009)*, Boulder, Colorado. June 2009.

PEER-REVIEWED WORKSHOP PUBLICATIONS

- 2015 Michael J. Paul, Mark Dredze, David A. Broniatowski, Nicholas Generous. **Worldwide influenza surveillance through Twitter.** *AAAI Workshop on the World Wide Web and Public Health Intelligence*, Austin, Texas. January 2015.
- Byron Wallace, Michael J. Paul, Noémie Elhadad. **What predicts media coverage of health science articles?** *AAAI Workshop on the World Wide Web and Public Health Intelligence*, Austin, Texas. January 2015.
- 2014 Shiliang Wang, Michael J. Paul, Mark Dredze. **Exploring health topics in Chinese social media: an analysis of Sina Weibo.** *AAAI Workshop on the World Wide Web and Public Health Intelligence*, Quebec City. July 2014.
- Mark Dredze, Renyuan Cheng, Michael Paul, David Broniatowski. **HealthTweets.org: a platform for public health surveillance using Twitter.** *AAAI Workshop on the World Wide Web and Public Health Intelligence*, Quebec City. July 2014.
- 2013 Michael J. Paul, Byron C. Wallace, Mark Dredze. **What affects patient (dis)satisfaction? Analyzing online doctor ratings with a joint topic-sentiment model.** *AAAI Workshop on Expanding the Boundaries of Health Informatics Using AI (HIAI)*, Bellevue, WA. July 2013.
- Mark Dredze, Michael J. Paul, Shane Bergsma, Hieu Tran. **Carmen: a Twitter geolocation system with applications to public health.** *AAAI Workshop on Expanding the Boundaries of Health Informatics Using AI (HIAI)*, Bellevue, WA. July 2013.
- 2012 Michael J. Paul and Mark Dredze. **Experimenting with drugs (and topic models): Multi-dimensional exploration of recreational drug discussions.** In submission to the *AAAI 2012 Fall Symposium on Information Retrieval and Knowledge Discovery in Biomedical Text*, Arlington, VA. November 2012.
- Alex Lamb, Michael J. Paul, Mark Dredze. **Investigating Twitter as a source for studying behavioral responses to epidemics.** In submission to the *AAAI 2012 Fall Symposium on Information Retrieval and Knowledge Discovery in Biomedical Text*, Arlington, VA. November 2012.
- Atul Nakhasi, Ralph J. Passarella, Sarah G. Bell, Michael J. Paul, Mark Dredze, Peter J. Pronovost. **Malpractice and malcontent: Analyzing medical complaints in Twitter.** In submission to the *AAAI 2012 Fall Symposium on Information Retrieval and Knowledge Discovery in Biomedical Text*, Arlington, VA. November 2012.

Ralph J. Passarella, Atul Nakhasi, Sarah G. Bell, Michael J. Paul, Peter J. Pronovost, Mark Dredze. **Twitter as a source for learning about patient safety events.** To appear in the *AMIA 2012 Annual Symposium (American Medical Informatics Association)*, Chicago, IL. November 2012.

William M. Darling, Michael J. Paul and Fei Song. **Unsupervised part-of-speech tagging in noisy and esoteric domains with a syntactic-semantic Bayesian HMM.** In the *EACL 2012 Workshop on Semantic Analysis in Social Media*, Avignon, France. April 2012.

2008 Michael Paul and Roxana Girju. **AIRTA: an automatic interdisciplinary research topic advisor.** *NSF-sponsored Symposium on Semantic Knowledge Discovery, Organization and Use - Demo session*, New York University. November 2008.

INVITED TALKS

3/02/2015	University of Maryland, Baltimore County <ul style="list-style-type: none"> Host: Dr. Tim Finin and Dr. Nilanjan Banerjee Title: Topic modeling with structured-priors for text-driven science 	Baltimore, MD
2/27/2015	University of Colorado Boulder <ul style="list-style-type: none"> Host: Dr. Leysia Palen Title: Topic modeling with structured-priors for text-driven science 	Boulder, CO
2/23/2015	Rice University <ul style="list-style-type: none"> Host: Dr. Lydia E. Kavraki and Dr. Chris Jermaine Title: Topic modeling with structured-priors for text-driven science 	Houston, TX
2/18/2015	University of California, Santa Barbara <ul style="list-style-type: none"> Host: Dr. Subhash Suri Title: Topic modeling with structured-priors for text-driven science 	Santa Barbara, CA
1/21/2015	National Cancer Institute (NIH) <ul style="list-style-type: none"> Host: Dr. Wen-ying Chou Title: Making sense of the web for public health using NLP 	Rockville, MD
12/03/2014	University of Maryland at College Park <ul style="list-style-type: none"> CLIP Colloquium Host: Dr. Jimmy Lin Title: Making sense of the web for public health using NLP 	College Park, MD
5/29/2014	Cafe Scientifique <ul style="list-style-type: none"> Monthly science talk series for the general public Host: Danielle Lucid Title: Using social media to study public health 	Annapolis, MD
9/16/2013	Association of Public Data Users <ul style="list-style-type: none"> Annual conference, Big Data session Host: Robert Scardamalia Title: Data mining social websites for public health 	Washington, DC
12/08/2010	University of Maryland at College Park <ul style="list-style-type: none"> CLIP Colloquium Host: Dr. Jordan Boyd-Graber Title: Summarizing contrastive viewpoints in opinionated text 	College Park, MD

PROFESSIONAL EXPERIENCE

- Summer 2013, Summer 2014 *Intern*
Microsoft Research Redmond, WA
Supervisors: Dr. Eric Horvitz, Dr. Ryen White
- Mined search query and browsing logs to understand web activity pertaining to medical information.
- October 2011 – Present *Advisory Board*
Sickweather LLC Baltimore, MD
<http://www.sickweather.com>
- Provides geographic surveillance of illnesses in the United States by monitoring social networking sites.
- Summer 2011 *Intern, Relevance Team*
Twitter Inc. San Francisco, CA
Supervisor: Dr. Alek Kolcz
- Worked on modeling user interests.
- Fall 2009 – Summer 2010 *Research Programmer*
University of Illinois, Dept of Computer Science Urbana, IL
Supervisor: Dr. Dan Roth
- Worked on the Learning Based Java modeling language and library.
 - Implemented various machine learning algorithms and helped develop a framework for learning with complex input and output.
 - Made efficiency improvements, increasing speed by up to 16x.
- Summer 2009 *Research Assistant*
University of Illinois, Department of Linguistics Urbana, IL
Advisor: Dr. Roxana Girju
- Developed probabilistic models of reciprocal social interactions.
 - Worked on a National Endowment for the Humanities (NEH) project on a family of endangered Peruvian languages.
- Summer 2008 *Research Fellow*
Univ. of Illinois, Data Sciences Summer Institute Urbana, IL
- My group created a search engine of UIUC courses, capable of displaying a course's prerequisite to/for map and automatically finding courses of similar content.

STUDENT ADVISING

- Fall 2013 – Current Shiliang Wang
- Masters student at JHU
 - Project: Studying health topics in Chinese social media
- Fall 2013 – Current Michael Smith
- Masters student at JHU
 - Project: Studying influenza awareness and media in social media
- Spring 2012 – Fall 2012 Alex Lamb
- Undergraduate at JHU (now at Amazon)
 - Project: Building a classification model of influenza in Twitter
- Spring 2012 – Fall 2012 Hieu Tran
- Undergraduate at JHU (now graduate student at MICA)
 - Project: Building a geolocation system for Twitter

ACADEMIC SERVICE

Co-chair	<ul style="list-style-type: none">• PSB 2016 (Pacific Symposium on Biocomputing) session on Social Media Mining for Public Health Monitoring and Surveillance<ul style="list-style-type: none">◦ January 4-8, 2016, Big Island of Hawaii
Organizer	<ul style="list-style-type: none">• Third Annual Mid-Atlantic Student Colloquium on Speech, Language and Learning (MASC-SLL)<ul style="list-style-type: none">◦ 10/12/2013, University of Maryland Baltimore County
Reviewer (Journal)	<ul style="list-style-type: none">• Computer science / NLP journals:<ul style="list-style-type: none">◦ Human-Computer Interaction (2015)◦ Transactions on Information Systems (TOIS) (2014)◦ Journal of Natural Language Engineering (JNLE) (2012, 2013)◦ Knowledge and Information Systems (KAIS) (2013)• Medical and general science journals:<ul style="list-style-type: none">◦ PLOS ONE (2014, 2015)◦ Journal of the American Medical Informatics Association (2014)◦ Health Informatics Journal (2014)◦ American Journal of Preventative Medicine (2014)◦ Nicotine & Tobacco Research (2013, 2014)
Reviewer / Program Committee	<ul style="list-style-type: none">• Computer science / NLP conferences:<ul style="list-style-type: none">◦ 2016: AAAI◦ 2015: ACL, NAACL, ICML, KDD, EMNLP◦ 2014: ACL, EACL, ICWSM◦ 2013: NAACL, *SEM◦ 2012: NAACL• ACL 2014 Workshop on Entity Attribute Prediction in Social Media• NIPS 2013 Workshop on Topic Models
Secondary Reviewer	<ul style="list-style-type: none">• Computer Science / NLP conferences:<ul style="list-style-type: none">◦ AAAI 2014, WWW 2014, NIPS 2011, NAACL 2010
Departmental Committees	<ul style="list-style-type: none">• Student admissions (JHU) (2015)• Student recruiting (JHU) (2011, 2012)